

CaseLinker: 5 Sources, 500 Cases, and Scaling Considerations

Mrinaal Ramachandran

Graduate Student, Department of Computer Science, University of Massachusetts Amherst

Independent Research | github.com/mrinaalr/CaseLinker | [Live Demo](#)

Technical Report Series #3 | April 2026

Abstract

CaseLinker is an open-source system for aggregating, extracting, and analyzing public ICAC case reports across jurisdictions. Report #3 scales the dataset to 500 cases across 5 sources — AZICAC, Georgia Bureau of Investigation, NCMEC CyberTipline, Idaho Office of Attorney General, and Michigan State Police — extracting 5,236 features across 9 dimensions via a hybrid deterministic-ML pipeline. This report introduces facet tree navigation for interactive case analysis, documents three primary data use cases (historical trend analysis, offender and crime evolution, and jurisdictional LE response shifts), and formalizes adversarial risk via a utility asymmetry proof showing that defender utility scales with dataset size while adversary utility remains approximately constant. Current coverage represents 5 public ICAC reports. Report #4 targets 2,000 cases and practitioner validation.

Keywords: ICAC investigations, cross-case analysis, scaling systems, historical trends, adversarial risk assessment, interpretable AI

Contents

1	Introduction	3
1.1	Monthly Report Series	3
1.2	ICAC Cases are Not Static, Variations in ICAC Operations	3
2	Improvements in Report #3	3
2.1	New Functions	3
2.2	Data Utility Analysis	4
3	Scaling Considerations	5
3.1	From 500 to Comprehensive Coverage	5
3.2	Challenges and Solutions	6
4	Data Use Cases	7
4.1	Broad Historical Comparisons and Evolution: 1998–2026	7
4.2	Offender and Crime Evolution	8
4.3	LE Response Shifts Across Jurisdictions	8
5	Facet Tree: 9-Dimensional Navigation	8
5.1	From Chaos to Precision	8
6	Adversarial Usage Risk Assessment	9
6.1	Setup and Assumptions	10
6.2	Proposition 1: Re-identification Risk is Source-Bounded	10
6.3	Proposition 2: Operational Signal is Source-Bounded	11
6.4	Proposition 3: Failure Conditions are Not Identifiable	11
6.5	Bounded Misuse Cases	12
6.6	Proposition 4: Utility Asymmetry	12
6.7	Access Controls and Risk Boundary Conditions	13
6.8	Summary of Results	14
7	Conclusion	14

1. Introduction

1.1. Monthly Report Series

Report #2 established hybrid deterministic-NER extraction, achieving 87.9% agency coverage across 207 cases (2011–2024), revealing distributed LE networks and early platform signals [2]. Report #3 scales to 500 cases across 3 new sources (Georgia Bureau of Investigation (GBI), Idaho Office of Attorney General, and Michigan State Police), introduces facet tree navigation, and paves the path for comprehensive historical and ICAC trend analysis over the program’s nearly three decades.

1.2. ICAC Cases are Not Static, Variations in ICAC Operations

An immediate conclusion from ingesting the GBI cases was the volume of large-scale, multi-perpetrator operations completed by this jurisdiction. A total of 56 known investigations were extracted from the 66 GBI case reports, 19 of which involved multiple perpetrators (34%), 4 of which were online, 1 proactive, and 14 undercover. Compare this against the AZICAC dataset, which shows only 2 cases involving multiple perpetrator operations from the 30 known operations (7%). Further analysis confirms that while each ICAC jurisdiction handles cases across all topics (CSAM, production, family, hands-on, grooming, platform involvement, investigation type), extracted patterns vary in distribution across the 9 dimensions.¹

2. Improvements in Report #3

2.1. New Functions

Facet Tree Navigation: The facet tree provides interactive, 9-dimensional navigation across the 500-case dataset, converting extracted case features into a precise, investigator-controlled tree. Dimensions include:

- (1) **Topic** (CSAM, possession, production, online-only)
- (2) **Severity** (infant victims, hands-on abuse, grooming, multiple perpetrators)

¹Preliminary conclusions directly from querying the CaseLinker interface. Their derivation from the facet tree is shown in 5.1. Underlying extraction accuracy and false positives have been checked via the audit system. However, until ablation tests are completed and extraction is rigorously tested, these stats are still an approximation over the dataset, not a hard claim. NHSR determination: HRPO #7668. Public, redacted data analyzed under UMass Amherst HRPO #7668 (Not Human Subjects Research). Refer to sources tab for full disclosure and compliance documentation.

- (3) **Platform** (Discord, Snapchat, Facebook, chat, social media)
- (4) **Inv. Type** (undercover, proactive, reactive, online)
- (5) **Source** (AZICAC, GBI, NCMEC, Michigan, Idaho)
- (6) **Agency** (from 518 unique agencies)
- (7) **Organization** (broader organizations, not necessarily law enforcement)
- (8) **Location** (NER extracted locations)
- (9) **Severity Phrase** (perpetrator admissions and case severity phrases)

New Sources: 66 cases from GBI, 244 cases from NCMEC CyberTipline annual report, 134 from Idaho Office of Attorney General and 11 from Michigan State Police.

2.2. Data Utility Analysis

This report introduces a comprehensive analysis of CaseLinker’s data utility across key use cases

- Historical trend analysis tracking platform/offender/crime evolution from 2010–2026
- ICAC responses and evolution (operation outcomes/case load) across sources/jurisdictions
- A formal risk argument, under stated assumptions, that CaseLinker adds no information beyond its public sources and that enforcement failure conditions are not identifiable from success-only records
- Bounded risk assessment (mosaic re-identification capped by the public source collection; evasion limited by absence of failure modes and operational details)

Full elaboration follows in Sections 4, 5, and 6 below.

Table 1: Source coverage (500 cases). “Features” = non-empty values among extracted dimensions.

Source	Cases	≥ 3 feat.	Inv. type	Severity	Platform
AZICAC	47	100.0%	63.8%	80.9%	27.7%
GBI	66	100.0%	84.8%	74.2%	98.5%
NCMEC	244	100.0%	59.2%	38.5%	33.9%
Idaho ICAC	134	100.0%	31.3%	35.8%	17.2%
Michigan ICAC	11	100.0%	100.0%	18.2%	100.0%

3. Scaling Considerations

3.1. From 500 to Comprehensive Coverage

There are 61 ICAC task forces, and the program has investigated and arrested offenders since 1998 [4]. No aggregate database or source collects the outcomes of ICAC or CSAM-related cases, so the total number of public ICAC case reports has to be estimated. The estimate below averages the per-year publication rates of the current task-force sources and scales that rate across all 61 task forces.

Table 2: Sample Sources for Public ICAC Case Volume Estimation. Years are counted inclusively.

Source	Date Range	Cases	Years	Cases/Year
AZICAC	2011–2014	47	4	11.75
GBI CEACC	2010–2026	66	17	3.88
Idaho AG ICAC	2021–2026	134	6	22.33
Michigan State Police	2025–2026	11	2	5.50
Average				10.87

Average annual publication rate: $\bar{r} = 10.87$ cases/source/year across 4 task-force sources. NCMEC is excluded from this estimate as a national aggregator rather than a single jurisdiction. Because 2026 is a partial year, rates for sources that extend into 2026 are slightly understated.

A naive linear scaling across all 61 task forces from 1998 through 2026 (29 years) gives

$$61 \times 10.87 \times 29 \approx 19,200.$$

This overstates the public record. Linear scaling assumes a constant publication rate, but ICAC case volume grew sharply after 2010, and public reporting from 1998–2009 is sparse. Restricting the same rate to 2010–2026 (17 years), the period in which public case reporting is substantial, gives

$$61 \times 10.87 \times 17 \approx 11,300.$$

This report therefore uses an estimate on the order of 10,000 public ICAC case reports, with the linear figure as an upper scenario.

Current dataset represents $500/10,000 = 5\%$ coverage. At scale, CaseLinker would analyze the complete census of successful public ICAC operations across all jurisdictions and agencies, enabling comprehensive trend analysis of online harms and their enforcement. This depends on whether public archives exist for each jurisdiction. A preliminary scan shows that most jurisdictions, including Texas, California, and Florida, publish public reports.

Ingestion and Processing: Case analysis is currently pre-computed; PDFs are ingested, processed, and analyzed with a one-time cost. On a local machine, the current end-to-end pipeline takes roughly 5 minutes on 500 cases to extract 5,236 features from 5 PDF sources, complete clustering, and store within the local SQLite / deployed PostgreSQL server. The projected time for 10,000 cases is 1.5–2 hours.

3.2. Challenges and Solutions

One emerging challenge is not computational resources and scale but rather issues relating to data bias, normalization, reporting details, and downstream impacts from differences in the length and quality of reporting. For example, the Michigan ICAC cases tend to be brief and succinct. This does *not* mean that there was not a substantial case load, investigative burden, or offender patterns worth studying, simply that the system was unable to extract it from the reported data as it was presented.

A quick visual example is comparing the cases reported by the AZICAC and Michigan

ICAC respectively. AZICAC includes detailed accounts of the perpetrator behaviors, violations, and even severity phrases indicating offender attributes *ex: azicac_2011_005*. Michigan ICAC cases appear to be higher level, summarizing charges, investigative methods and agencies but lack explicit details on the offenders. This disparity continues within reports, cases within the NCMEC 2024 annual report vary, with some having brief success outcomes and prosecutions and others having detailed insights into the specific crime.

These are limitations at the available data level, not analytical level. These are also important limitations to have, as public reports have mosaic reidentification and adversarial risks, which are discussed in section 6. While there are different methods for normalization, at this time *no effort will be made to normalize this data*. To be clear, all cases will be extracted, processed, and analyzed in their current format, regardless of reporting detail. This is to serve CaseLinker’s primary functionality: visualizing and analyzing public reports. Downstream statistical analysis and broader inferences would require normalization, especially if attempting to come to conclusions on platform involvement for example.

As a final footnote, the analytical layer can be adapted. The current system supports broad feature extraction across the 9 dimensions to collect as much information as possible for investigators, journalists, and case studies. However, the processing layer is not set in stone and can be adapted to be focused on statistics, though that is not the primary research direction CaseLinker’s public release is evolving towards.

4. Data Use Cases

4.1. Broad Historical Comparisons and Evolution: 1998–2026

- Cross-era platform involvement: early internet → P2P → social media → modern platforms, traceable across the full date range of the dataset
- Case volume and charge type shifts across three decades, filterable by source, jurisdiction, and year range
- Correlation of caseload surges with legislative changes, major platform policy shifts, and new technology.

4.2. Offender and Crime Evolution

- Distribution of offense type (possession, production, distribution, hands-on) over time and across jurisdictions
- Shifts in grooming behavior, online-only vs. hands-on operations, and multi-perpetrator coordination across the dataset
- Perpetrator admission patterns and severity phrases extracted across 500 cases, enabling longitudinal behavioral comparison

4.3. LE Response Shifts Across Jurisdictions

- Investigation type distribution (undercover, proactive, reactive) by source and agency, revealing how different task forces allocate investigative effort
- Multi-perpetrator operation rates vary significantly by jurisdiction: GBI at 34% vs. AZICAC at 7%, invisible without cross-source aggregation
- Recurring agency networks and inter-agency coordination patterns across the 518 unique agencies in the dataset

5. Facet Tree: 9-Dimensional Navigation

5.1. From Chaos to Precision

The initial tree represents 500 distinct cases and contains 2,459 nodes. This immediately demonstrates that there are a variety of branches that can be taken to group, analyze, partition, and otherwise engage with the extracted features from these cases.

The method presented is a facet tree with interactive pruning and drill-down to dynamically focus on topics of interest.

Analysts can select one or more topics from each dimension to analyze, both individually and in relation to other dimensions. The conclusions on GBI responses in this report were reached through the following navigation steps:

1. Select Split on: *Source* and narrow to GBI, pruning the four other sources from the tree, reducing total cases from 500 to 66 and nodes from 2,459 to 298.

2. Identify topics of interest. To support analysis of investigation types, select all known investigation types, pruning the tree further to include 56 cases and 256 nodes.
3. To case study GBI larger operations involving multiple perpetrators (large scale operations likely involving weeks–months of planning and resource allocation), select the multiple perpetrators tag from severity indicators. This narrows the search to 19 cases and 93 nodes.
4. Remove all splits that do not improve knowledge for studying GBI operations involving multiple perpetrators (severity phrases, location, organization, agency, topic, platform). This reduces nodes from 93 to 11.
5. The resulting tree (referenced below) shows GBI cases involving multiple perpetrators, split by investigation type. Analysts could instead split by platform, victim age, date range, location, or any other combination and permutation of interest.




Figure 1: Resulting Facet tree from running the pruning steps above.

6. Adversarial Usage Risk Assessment

Below is a formal analysis of the adversarial risk of the CaseLinker system. As argued in my position piece, Why Internet Crimes Against Children And Retrospective Case Analysis Matters [3], all technology can be misused by bad actors and CaseLinker is no

exception. The question is not whether CaseLinker presents *no risk*, but whether defender utility outweighs adversary insight and whether appropriate steps are taken to prevent misuse. The results below hold under explicitly stated assumptions, and each result is scoped to what those assumptions support.

6.1. Setup and Assumptions

Let $\mathcal{S} = \{s_1, s_2, \dots, s_m\}$ denote the collection of public source documents (law enforcement press releases and annual reports), and let $\mathcal{D} = f(\mathcal{S}) = \{c_1, c_2, \dots, c_n\}$ denote the CaseLinker dataset of n processed case records, where f is the full CaseLinker pipeline (extraction, NER, clustering, facet tree, and aggregate statistics). Let X denote any quantity an adversary may wish to learn, and let $I(X; \cdot)$ denote mutual information: the reduction in uncertainty about X , in bits, obtained by observing the argument.

- **A1 (Source closure).** f depends on X only through $\mathcal{S}$. Every feature in $\mathcal{D}$ is computed from the public documents; components trained on external data (e.g., NER models) use no case-specific information beyond $\mathcal{S}$. Formally, $X \rightarrow \mathcal{S} \rightarrow \mathcal{D}$ is a Markov chain.
- **A2 (Publication policy).** The sources report successful outcomes (arrests, charges, convictions). Agencies do not publish failure conditions, tip routing, undercover persona construction, or forensic tooling details.
- **A3 (Public access).** Any adversary can already access $\mathcal{S}$; the documents are published for public release.

The analysis rests on the data processing inequality [6]: for any Markov chain $X \rightarrow \mathcal{S} \rightarrow f(\mathcal{S})$,

$$I(X; f(\mathcal{S})) \leq I(X; \mathcal{S}).$$

Processing can preserve or discard information about X , but cannot create it.

6.2. Proposition 1: Re-identification Risk is Source-Bounded

Proposition 1. Let X be the identity of any individual named or described in the sources. Under A1, $I(X; \mathcal{D}) \leq I(X; \mathcal{S})$.

Argument. Immediate from the data processing inequality under A1. CaseLinker does not process victim names, perpetrator personal identifiers, addresses, or case numbers beyond what appears in structural metadata, so in practice the processed features retain

less identifying information than the sources, not more. The re-identification ceiling is therefore set by what the originating agencies chose to publish and redact. The bound is stated over the full source collection $\mathcal{S}$, not a single document: linking two public releases about the same case can reveal more than either release alone, but that linkage is already available to anyone reading the public record (A3), and CaseLinker cannot exceed it. $\square$

6.3. Proposition 2: Operational Signal is Source-Bounded

Proposition 2. Let Θ denote law enforcement operational variables (interception methods, undercover persona construction, tip source routing, forensic tooling). Under A1, $I(\Theta; \mathcal{D}) \leq I(\Theta; \mathcal{S})$, for every n .

Argument. Immediate from the data processing inequality under A1. The bound does not claim that the sources contain no operational information: press releases carry investigation type labels (undercover, proactive, reactive), platform mentions, agency names, charge categories, and outcomes. The claim is that whatever operational signal exists is already in the public record, and aggregation does not recover what was never published. Under A2, the operational detail most useful for evasion is absent from $\mathcal{S}$, and therefore from $\mathcal{D}$. $\square$

6.4. Proposition 3: Failure Conditions are Not Identifiable

Proposition 3. Let $Y \in \{\text{success}, \text{failure}\}$ denote the outcome of an enforcement effort and x the observable features of a case. Under A2, $P(Y = \text{failure} \mid x)$ is not identifiable from $\mathcal{S}$ or $\mathcal{D}$, for any n .

Argument. Under A2, every record satisfies $Y = \text{success}$, so the data are draws from $P(x \mid Y = \text{success})$ alone. By Bayes' rule,

$$P(Y = \text{failure} \mid x) = \frac{P(x \mid \text{failure}) P(\text{failure})}{P(x)},$$

which depends on $P(x \mid \text{failure})$ and $P(\text{failure})$, neither of which appears in the data. Any two worlds that share $P(x \mid \text{success})$ but differ in $P(x \mid \text{failure})$ generate identically distributed records, so no estimator, at any sample size, can distinguish them. An adversary seeking the conditions under which enforcement fails requires failure data; neither $\mathcal{S}$ nor $\mathcal{D}$ contains any. $\square$

A practical consequence: the absence or rarity of a feature in $\mathcal{D}$ (a platform, a jurisdiction, a victim age group) is ambiguous. Rarity can reflect low enforcement, low offending, low

platform usage, or reporting practice, and Proposition 3 implies the data cannot separate these explanations.

6.5. Bounded Misuse Cases

Despite the above, three adversarial misuse vectors are worth naming explicitly, along with their practical ceiling.

Platform Avoidance. An adversary observing that Facebook appears in a repeated number of successful prosecutions may attempt to migrate to less-documented platforms. This signal is available from repeated manual reads of any public ICAC source, and by Proposition 3 a platform’s rarity in $\mathcal{D}$ does not indicate weaker enforcement on that platform. Platform avoidance is further limited by CyberTipline reporting and cross-platform device correlation, which can bind an offender to a device independent of the platform used.

Geographic and Agency Avoidance. An adversary may attempt to infer which jurisdictions or agencies produce high prosecution rates and avoid them. However, internet-based ICAC offenses are not geographically bounded in the same way as physical crimes: CyberTipline reports and device forensics route investigations to the offender’s location regardless of which jurisdiction initiates the case. Geographic avoidance provides minimal protective value.

Victim Demographic Targeting. An adversary may attempt to use severity and victim age distributions to identify demographic groups that appear less prosecuted. Proposition 3 applies directly: a group’s lower frequency in a dataset of successful prosecutions cannot be distinguished from lower offending against that group or from reporting practice, so the distributions do not reveal enforcement blind spots. CaseLinker surfaces no information on how to obtain CSAM, how offender communities form, or how platforms are regulated beyond what is publicly known.

6.6. Proposition 4: Utility Asymmetry

The data processing inequality applies to defenders as well: CaseLinker creates no new information for anyone. Its value lies in cost. Computing a cross-case statistic from $\mathcal{S}$ by hand requires reading every relevant document; computing the same statistic from $\mathcal{D}$ requires a query. The asymmetry concerns *which* quantities become cheap.

Proposition 4. Under A1–A3, CaseLinker lowers the cost of computing distributional

summaries of successful cases for defenders and adversaries alike, but the evasion-relevant quantities Θ and $P(Y = \text{failure} \mid x)$ gain nothing from this cost reduction: Θ is bounded by its availability in the public record (Proposition 2), and failure conditions are not identifiable at any cost or sample size (Proposition 3).

Argument. The defender questions CaseLinker answers (cross-jurisdictional success patterns, platform evolution, multi-perpetrator operation typologies, recurring agency networks, offense type distributions) are summaries of $P(x \mid \text{success})$, and their precision grows with n . For an estimated proportion $\hat{p}$ from n cases, the standard error is $\sqrt{\hat{p}(1 - \hat{p})/n}$: the GBI multi-perpetrator rate of $19/56 \approx 34\%$ carries a standard error of roughly 6 percentage points, and comparisons across jurisdictions sharpen as coverage grows. These summaries are expensive to assemble manually and cheap to query, which is where defender utility comes from. The adversary’s evasion problem requires Θ or $P(Y = \text{failure} \mid x)$; Propositions 2 and 3 bound both independently of n . The adversary does gain the same cheap access to success-pattern summaries (for example, platform frequency tables), and the Bounded Misuse Cases above argue that this access has low evasion value. $\square$

6.7. Access Controls and Risk Boundary Conditions

The above propositions hold for the current system and data scope. Two boundary conditions warrant explicit statement.

Access Control. For cases involving elevated sensitivity indicators (infant victims, hands-on severity phrases), CaseLinker restricts direct case retrieval to access-key holders. Facet navigation and aggregate statistics remain public. This prevents the system from functioning as a targeted lookup tool for the most sensitive case details while preserving full analytical utility for investigators, researchers, and the general public.

Scope Boundary. This risk assessment is bounded to the current data scope: public, redacted case summaries. If the system were extended to incorporate non-public case data, active investigation records, or unpublished operational details, A1–A3 would no longer hold and the threat model would require fundamental reassessment. Similarly, if future extraction surfaces qualitatively new sensitive features not present in the current 9-dimensional schema, risk must be re-evaluated. The bounds also concern $\mathcal{D}$ alone: an adversary combining CaseLinker output with external non-public data falls outside this model. The propositions above do not generalize beyond public redacted sources.

6.8. Summary of Results

Public ICAC reports record successful operations only. Adversary gains: cheaper access to success-pattern summaries such as “avoid platform x ” (available from manual reads and weak for evasion). Defender gains: cross-case success patterns across 61 task forces since 1998.

Re-identification: Bounded by the public source collection (Proposition 1). No names are processed.

Operational signal: Bounded by what agencies published (Proposition 2).

Failure conditions: Not identifiable from success-only records at any n (Proposition 3).

Misuse Ceiling: Platform and geographic avoidance limited by CyberTipline reporting and device forensics. Apparent demographic gaps cannot be separated from offending rates or reporting practice.

7. Conclusion

Report #3 scales CaseLinker to 500 cases across 5 sources, introduces facet tree navigation across 9 dimensions, and formalizes the adversarial risk argument that has been implicit since the project began. The release extracts 5,236 features, identifies 518 unique agencies, and builds a facet tree of 2,459 nodes.

Those numbers serve a larger purpose: ICAC enforcement is a 28-year-old ecosystem that has lacked a public tool for looking at itself in aggregate. Jurisdictions operate in parallel, catching offenders and publishing successful outcomes, but lack mechanisms for researchers, journalists, or task force commanders to ask: what does the full picture look like? Which platforms recur? How do multi-perpetrator operations distribute across jurisdictions? What do perpetrator admissions look like at scale? Which funding, policy, and enforcement actions should be taken based on emerging trends? These are not easy questions to answer. However, they are the baseline questions any serious analysis of a law enforcement ecosystem should start with, and until now they have been unanswerable without months of manual work and reviewing hundreds of graphic cases across thousands of pages.

Five percent coverage of estimated public ICAC reports is the beginning. The analytical value compounds with scale: Proposition 4 shows that the precision of cross-case patterns grows with coverage, while the adversary’s evasion problem does not get easier. At 500 cases, you have a window into this ecosystem. At 10,000 you can see the structure. That is the direction this project is going.

CaseLinker is not a replacement for investigative judgment, institutional knowledge, or the expertise of the people who have been doing this work for decades. Often, investigators, commanders, and experts who have been protecting children, catching offenders, and shaping policy recognize patterns and trends extremely well. CaseLinker is a way of making what is known visible at scale, uncovering deeper patterns that manual review cannot reach, supporting these claims in an auditable, concrete, evidence-tied manner, and providing material insights into the nature of these crimes for those unaware, without requiring the consumption of graphic case material. The goal has not changed since the initial release: give investigators, researchers, and the general public the ability to understand this landscape without requiring repeated direct exposure to its worst material [1, 5].

Report #4 targets 2,000 cases, semantic similarity across cases, and the first round of practitioner validation.

CaseLinker is available at github.com/mrinaalr/CaseLinker under the MIT License.

Live demo: web-production-13a2.up.railway.app.

References

- [1] Perez, L. M., Jones, J., Englert, D. R., & Sachau, D. (2010). Secondary traumatic stress and burnout among law enforcement investigators exposed to disturbing media images. *Journal of Police and Criminal Psychology*, 25(2), 113–124.
- [2] Ramachandran, M. (2026, March). *CaseLinker: Interpretable ML Approaches for Analyzing Internet Crimes Against Children Reports. Technical Report Series #2*. Uni-

- versity of Massachusetts Amherst. <https://mrinaalr.github.io/website/CaseLinker-%20Interpretable%20ML%20Approaches%20for%20Analyzing%20Internet%20Crimes%20Against%20Children%20Reports.pdf>
- [3] Ramachandran, M. (2026, March). *Why Internet Crimes Against Children And Retrospective Case Analysis Matters*. University of Massachusetts Amherst. <https://mrinaalr.github.io/website/Why%20Internet%20Crimes%20Against%20Children%20And%20Retrospective%20Case%20Analysis%20Matters.pdf>
- [4] Office of Juvenile Justice and Delinquency Prevention. (n.d.). *Internet Crimes Against Children Task Force Program*. Retrieved April 2026. <https://ojjdp.ojp.gov/programs/internet-crimes-against-children-task-force-program>
- [5] Burns, C. M., Morley, J., Bradshaw, R., & Domene, J. (2008). The emotional impact on and coping strategies employed by police teams investigating internet child exploitation. *Traumatology*, 14(2), 20–31.
- [6] Cover, T. M., & Thomas, J. A. (2006). *Elements of Information Theory* (2nd ed.). Wiley. Theorem 2.8.1 (data processing inequality).